

Segmentation and Paging Algorithms

44-550: Operating Systems

Advantages

- No external fragmentation
- Simple memory management algorithm
- Swapping is easy
- Transparent to programmer (system allocates memory)

Disadvantages

- Internal fragmentation on last page of process memory
- Page tables may consume more memory
- Doesn't really offer memory protection
 - Can only flag whether or not a frame is valid or invalid

Segmentation

- Variable length modules of a program corresponding to logical units (addresses)
- Represent modular structure of a program
- Examples include:
 - Main method
 - Other methods
 - Data
- We *prefer* to load segments into contiguous blocks of memory
- Very common way to achieve memory protection
- Segments allocated in various sizes depending on the process' need for address space
- Note: At this point, if you google paging vs segmentation, the results you get will be hairy; Much of the difference depends on what level we're talking "segmentation". Additionally, some Windows documentation refers to swapping as paging...

Segmentation

- The OS keeps track of segments for a process including starting address and length.
- Linux avoids low level segmentation as much as possible:
 - Kernel Code
 - Kernel Data
 - Userspace Code
 - Userspace Data
- But wait... what's a segmentation fault?
 - Program tried to access an invalid or illegal memory address (like NULL) or a value larger than a valid pointer for that process
 - Not really related to the segmented memory model.
 - Different from a *Page fault*
- Operating systems can overlay segmentation on top of paging: a process has segments, and those segments directly map to one or more pages

Virtual Memory

- Memory space of a process is divided into blocks of pages or segments
 - Doesn't matter what, for our purposes
- Not all blocks are required to execute a process
- Keep entire program address space on disk
 - BUT, we usually have a lot less physical RAM than we do disk space.
- The virtual memory manager swaps out pages/segments of an executing process as it's needed
- The OS must provide a way to translate between virtual and physical addresses
- Virtual address space is MUCH larger than physical address space
 - Otherwise, a 30GB game (not unheard of) wouldn't fit in 4-8GB of ram
- An efficient way to load blocks as the program executes is needed
 - Either swap space or page file

Process Locality and Locality of Reference

- The set of pages a process operates on are called a locality
- Locality of Reference is the concept that data accesses are predictable; can be exploited to have an intelligent virtual memory manager
- Temporal Locality
 - Data recently used is likely to be reused
- Spatial Locality
 - Data element X was used, so data close to it in physical memory may be used soon
- Exploited to avoid *page faults* as much as possible
 - If required page is not in memory, must go to swap space to fetch it
 - This is a slow process, since we are accessing a slower kind of memory (HDD/SDD)

Address Translation

- ① Check for valid page number
- ② Check for valid physical address
 - If page is not in memory, OS must trigger an interrupt to access disk (page fault)
- ③ OS must write data pages back to disk if there are modifications to physical memory (must keep the page/swap file up to date)
 - Hooray consistency!
- Managing the page table can be costly; you may only need to load part of it into memory
- Page size is important too
 - Too big: taking time to load spurious data
 - Too small: Lots of page faults

Paging with Virtual Memory

- Large processes may require multiple pages
- Page faults occur when the process references a page not in memory
 - This is a very costly operation

- Fetch policy: when to swap a page into memory
- Replacement policy: If there are no empty frames, which page gets replaced?
- Placement policy: where to place the page in memory
- Number of frames to allocate to a process

- Demand paging: page isn't loaded until referenced in memory
 - Page faults occur when an instruction or an operand to an instruction is needed
- Prepaging: Load page before it is referenced by the program
 - Must be careful about which pages to load (what happens if they don't get used soon?)

- When a page fault occurs:
 - Process generating fault is suspended
 - OS locates referenced page on disk using page tables
- if there is no free frame, a page must be selected to be moved back to the HDD
- Load the referenced page into the selected frame and update page/frame tables
- Resume interrupted process
- If frame sent back to memory has been modified, need to update the page in memory

- How many frames should a process be allowed?
- Too few frames means lots of page faults
 - Causes *thrashing*: worse than deadlock because nobody makes progress, and the system can be brought down
 - Also: excessive thrashing could (possibly) reduce the life of SSDs (though not noticeably)
- Can use equal or proportional allocation (proportional to size)
- Replacement can use either local or global allocation
 - local: select only from process' frames
 - global: select from other process frames or empty ones

- Time to service page fault is important (lots of overhead)
 - Time interval to service page fault interrupt
 - Time to swap out replaced page to disk
 - Time to swap in (load) referenced page from disk
 - Delay in queueing for HDDs
 - Delay in scheduling process with referenced page
- We can investigate how much this hurts by guessing at the page fault rate ($0 \leq p \leq 1.0$):
 - If $p = 0$, no page faults
 - If $p = 1$, every reference is a fault

Effective Access Time (EAT)

Effective Access Time

If we assume that checking to see if the page we want is already loaded is instantaneous (0 time): $EAT = (1 - p)(MAT) + p(FST)$

MAT: memory access time; FST: Fault Service Time

Tells us on average how long memory accesses take

Example: $MAT = 200ns$, $FST = 8ms = 8000000ns$. Then:

$$\begin{aligned} EAT(p) &= 200(1 - p) + 8000000p \\ &= 200 - 200p + 8000000p \\ &= 200 + 7999800p \end{aligned}$$

p	$EAT(p)$
0	.2us
$\frac{1}{1000}$	8.2us(40x slower)
1	8ms(40000x slower)

- Assume a process has a stream of pages it will reference (per process)
- Assume also that each action is represented by one page (for simplicity)
 - Example: $\langle 3, 6, 1, 3, 3, 2, 7, 8, 9, 10, 11 \rangle$
- We have 3 static policies we can use:
 - FIFO
 - Optimal replacement
 - LRU

- As name suggests, the first page to get replaced will be the first to be swapped out
- Our example will use a 3 page “segment”
 - All this means is that we have 3 pages that can be loaded at a time
- Note: This can exhibit poor performance even if the number of frames is increased

Needed Page:	4	7	3	0	1	7	3	8	5	4	5	3	4	7
Page 1:														
Page 2:														
Page 3:														
Hit/Fault														

FIFO Example

Needed Page:	4	7	3	0	1	7	3	8	5	4	5	3	4	7
Page 1:	4													
Page 2:														
Page 3:														
Hit/Fault	F													
Needed Page:	4	7	3	0	1	7	3	8	5	4	5	3	4	7
Page 1:	4	4												
Page 2:		7												
Page 3:														
Hit/Fault	F	F												

FIFO Example

Needed Page:	4	7	3	0	1	7	3	8	5	4	5	3	4	7
Page 1:	4	4	4											
Page 2:		7	7											
Page 3:			3											
Hit/Fault	F	F	F											

Needed Page:	4	7	3	0	1	7	3	8	5	4	5	3	4	7
Page 1:	4	4	4	0										
Page 2:		7	7	7										
Page 3:			3	3										
Hit/Fault	F	F	F	F										

FIFO Example

Needed Page:	4	7	3	0	1	7	3	8	5	4	5	3	4	7
Page 1:	4	4	4	0	0									
Page 2:		7	7	7	1									
Page 3:			3	3	3									
Hit/Fault	F	F	F	F	F									

Needed Page:	4	7	3	0	1	7	3	8	5	4	5	3	4	7
Page 1:	4	4	4	0	0	0								
Page 2:		7	7	7	1	1								
Page 3:			3	3	3	7								
Hit/Fault	F	F	F	F	F	F								

FIFO Example

Needed Page:	4	7	3	0	1	7	3	8	5	4	5	3	4	7
Page 1:	4	4	4	0	0	0	3							
Page 2:		7	7	7	1	1	1							
Page 3:			3	3	3	7	7							
Hit/Fault	F	F	F	F	F	F	F							

Needed Page:	4	7	3	0	1	7	3	8	5	4	5	3	4	7
Page 1:	4	4	4	0	0	0	3	3						
Page 2:		7	7	7	1	1	1	8						
Page 3:			3	3	3	7	7	7						
Hit/Fault	F	F	F	F	F	F	F	F						

FIFO Example

Needed Page:	4	7	3	0	1	7	3	8	5	4	5	3	4	7
Page 1:	4	4	4	0	0	0	3	3	3					
Page 2:		7	7	7	1	1	1	8	8					
Page 3:			3	3	3	7	7	7	5					
Hit/Fault	F	F	F	F	F	F	F	F	F					

Needed Page:	4	7	3	0	1	7	3	8	5	4	5	3	4	7
Page 1:	4	4	4	0	0	0	3	3	3	4				
Page 2:		7	7	7	1	1	1	8	8	8				
Page 3:			3	3	3	7	7	7	5	5				
Hit/Fault	F	F	F	F	F	F	F	F	F	F				

FIFO Example

Needed Page:	4	7	3	0	1	7	3	8	5	4	5	3	4	7
Page 1:	4	4	4	0	0	0	3	3	3	4	4			
Page 2:		7	7	7	1	1	1	8	8	8	8			
Page 3:			3	3	3	7	7	7	5	5	5			
Hit/Fault	F	F	F	F	F	F	F	F	F	F	H			
Needed Page:	4	7	3	0	1	7	3	8	5	4	5	3	4	7
Page 1:	4	4	4	0	0	0	3	3	3	4	4	4		
Page 2:		7	7	7	1	1	1	8	8	8	8	3		
Page 3:			3	3	3	7	7	7	5	5	5	5		
Hit/Fault	F	F	F	F	F	F	F	F	F	F	H	F		

FIFO Example

Needed Page:	4	7	3	0	1	7	3	8	5	4	5	3	4	7
Page 1:	4	4	4	0	0	0	3	3	3	4	4	4	4	
Page 2:		7	7	7	1	1	1	8	8	8	8	3	3	
Page 3:			3	3	3	7	7	7	5	5	5	5	5	
Hit/Fault	F	F	F	F	F	F	F	F	F	F	H	F	H	

Needed Page:	4	7	3	0	1	7	3	8	5	4	5	3	4	7
Page 1:	4	4	4	0	0	0	3	3	3	4	4	4	4	4
Page 2:		7	7	7	1	1	1	8	8	8	8	3	3	3
Page 3:			3	3	3	7	7	7	5	5	5	5	5	7
Hit/Fault	F	F	F	F	F	F	F	F	F	F	H	F	H	F

$$p = \frac{12}{14} = .85714$$

$$EAT(.857) = 6.85ms$$

‘ Ouch

- Optimal:
 - Replace page in memory that will not be used for the longest period
 - Requires the entire page reference stream in advance (what's with all the psychic algorithms?)
- Least Recently Used:
 - Replace the page in memory that has not been used for the longest period
 - Exploits Locality of Reference (temporal)
- In class examples: next time