

CPU Scheduling, Part Deux

44-550: Operating Systems


- FCFS
 - Fair
 - Simple to implement
 - Far from optimal
- SJF
 - Not Fair
 - Requires ability to estimate CPU burst
 - Could push longer running jobs to never run
- What happens if we try a more sophisticated approach?

Round Robin (Yum!)

- Each process is given quanta (equal time slices)
- Takes turns of one quantum each
- Let's see what happens with the original example

Process	Burst Time(t)	Arrival Time
A	7	0
B	4	4
C	1	5
D	1	9
E	3	12

RR Example




Process	Required Time (T)	Wait Time	Total Time
A	7	5	12
B	4	6	10
C	1	1	2
D	1	1	2
E	3	2	5
Average		2.6	5.8
		Wait Time	Turnaround

- Benefits short burst processes greatly
- Can hurt the runtime of processes with long bursts
 - Doesn't excessively penalize them; they will not starve
- Preemptive

Shortest Remaining Time (SRT)

- When a process arrives in the ready queue with an expected CPU burst time smaller than the remaining time of the running processes, the new one preempts the running process



SRT Properties

- Long processes can starve
- Short processes get serviced quickly
- Works well in many cases
- Works poorly if estimation of remaining time is bad

A Real Life Implementation: CFQ

- Default scheduler for the Linux kernel since the 2.6.23 release
- Not based on run queues
 - A red black tree implements a timeline of future task execution
- Uses nanosecond granularity
 - We don't need no stinkin' time slices
- Does not need heuristics to determine the interactivity of a process
 - how often it hits up I/O

Red Black Trees

- Built on top of BST (for our example)
- A node is either red or black
- All leaves are black, and have the value NULL/NIL
- If a node is red, all children are black
- Every path from a node to any descendant NIL nodes has the same number of black nodes in the path
 - The uniform number of black nodes in the path from root to leaves is called the *black height* of the tree




Figure: From https://en.wikipedia.org/wiki/Red%E2%80%93black_tree

The Algorithm (as per wikipedia)

Processes are inserted into the tree based on their current spent execution time

- The left most node of the scheduling tree is chosen and sent for execution
 - it has the lowest spent execution time
- If the process completes, removed from the system and tree
- If the process reaches its maximum execution time or otherwise stopped, reinserted into the scheduling tree based on its new spent execution time.
- The new left most node will then be selected from the tree. Repeat

- Not Completely Fair (as the name implies)
- Processes that sleep a lot are not penalized (as their spent execution time remains low), and will be given priority when they need the CPU

Dynamic Scheduling Policies

- Processes can alternate between I/O and CPU bursts
 - We could give processes that need I/O followed by CPU higher priority to keep the I/O busy (bottleneck)
- How do we decide what processes to give high priority?
- Dynamic Scheduling adjusts the priority as the process runs
 - Predicts what will happen next (branch prediction is hard)
- We want to keep all the resources busy (I/O and CPU)

Possible Dynamic Scheduling Heuristic

- Allocate CPU to highest priority process
- When a process is selected to be executed, give it a time slice
- If a process makes an I/O request before it's timeslice is up, raise it's priority (assume it will make an I/O request shortly after a small CPU boost)
- If the process reaches the end of its timeslice without making an I/O request, lower the priority

Multiprocessor Scheduling

- So far, assumed only one processor
- Most computers this day have multiple cores (or multiple processors)
- Many different ways to accomplish scheduling across the cores
 - Have a scheduler data structure per core
 - Have one scheduler data structure, and distribute tasks from it to each core
 - ...