

Processes and Threads

44-550: Operating Systems

Processes and Threads

- Exist at execution time
- Have fast state changes
 - in memory
 - waiting
- A process:
 - is a fundamental computational unit
 - can have one or more threads
 - is handled by the process management module
 - requires **system** resources

- Process (job): program in execution, ready to execute, or waiting for execution
- A program is static; processes are dynamic
- Different types of processes:
 - user processes
 - system processes
- Different wait queues exist for different types of processes

- Major function of the OS
- OS manages which processes get what CPU/memory resources at what times
- This is called *multiprogramming*

- CPUs must be shared
- Scheduling minimizes idle time
- In Windows, the Task Manager will show how many processes are running
 - How many are there?
 - How many CPU cores are on your computer?
- On a Mac or Linux machine, run `top` or `htop` at the command line

- Big computer sciency term
- Abstraction is used with processes
- Every process gets:
 - Process Identifier (PID)
 - address
 - memory space
 - program code
 - data
 - resources required

Process Example


Ending the Infinite Background Loop

- Check out `top`
- Run `ps -fu <username>`
- Find the PID of the process
- Kill it

```
kill <pid>
```

```
kill -9 <pid> # hard kill
```


Process States and Their Allowable Transitions



- Run \rightarrow ready: interrupt
- Ready \rightarrow run: preempt or timeslice end
- Run \rightarrow wait/blocked: requesting unavailable resource
- Wait/blocked \rightarrow ready: getting resources

PCB: Process Control Block

- Also known as a process descriptor
- Created with each new process
- Contains all data associated with a process
- OS Queues use a reference (pointer) to a PCB so the queue doesn't need to contain the entire PCB

| | |
|---------------------------|-----|
| name | PID |
| process owner/user | |
| state | |
| list of threads | |
| list of resources | |
| list of child process | |
| address space | |
| privileges or permissions | |
| CPU register image | |

| |
|----------------------------------|
| PID |
| state (ready, run, wait/blocked) |
| CPU Registers |
| List of resources |
| List of child resources |
| Parent pointer |
| Permissions |
| Stack and code pointers |

- Are like “lightweight” processes
 - processes can contain many threads
 - threads change (dynamic within process)
- Current OSes support multithreading
 - multiple threads/tasks per process
- Multiple threads more efficient than multiple processes

Thread Descriptor

| |
|--|
| Thread ID |
| Context (program counter within process) |
| Execution Stack |
| Local memory |
| Reference to parent process (for shared resources) |
| Execution state |
| List of related threads |
| Thread Priority |
| Thread specific resources |


- Every process creates one thread on process creation
 - “main”
- Other threads spawned from the “main” thread
- Threads have same states as processes
- A process will terminate when all threads in the process terminate

Many Threads, One Process

```
#include <stdio.h>

int globvar1;
int globvar2;


int main(int argc, char* argv[])
{
    int mainvar;
    // do some cool stuff
    return 0;
}
```



Many Threads, One Process

```
#include <stdio.h>

int globvar1;
int globvar2;
// THIS PROTOTYPE IS INCORRECT
void foo(){
    int thvar;
    // do stuff
}
int main(int argc, char* argv[]){
    int mainvar;
    // do some cool stuff
    // launch thread calling foo
    return 0;
}
```



- Multiple threads may exist in one process; regardless the process has only one process descriptor
- Threads have less overhead than processes
- Operations
 - Create thread
 - terminate thread
 - switch between threads (context switching)
 - communication between threads
- All operations have a cost

- Threads share the code and resources of the parent process
 - No need to switch out code and resources when switching threads
- Thus, fewer cache misses!

More Threading!

- Threads exist at both user and kernel level
- User level (no kernel level intervention):
 - Windows
 - POSIX
 - Java...
- Kernel threads perform management tasks done by kernel processes
- A process can still be running even if a thread is blocked