

LRU, Optimal, and Caching

44-550: Operating Systems

Least Recently Used

- Replace the page in memory that has not been used for the longest amount of time

Needed Page:	0	1	2	1	4	2	3	7
Page 1:								
Page 2:								
Page 3:								
Hit/Fault								

LRU Example

Needed Page:	0	1	2	1	4	2	3	7
Page 1:	0*							
Page 2:								
Page 3:								
Hit/Fault	F							

Needed Page:	0	1	2	1	4	2	3	7
Page 1:	0*	0*						
Page 2:		1						
Page 3:								
Hit/Fault	F	F						

LRU Example

Needed Page:	0	1	2	1	4	2	3	7
Page 1:	0*	0*	0*					
Page 2:		1	1					
Page 3:			2					
Hit/Fault	F	F	F					

LRU Example

Needed Page:	0	1	2	1	4	2	3	7
Page 1:	0*	0*	0*	0*				
Page 2:		1	1	1				
Page 3:			2	2				
Hit/Fault	F	F	F	H				

Needed Page:	0	1	2	1	4	2	3	7
Page 1:	0*	0*	0*	0*	4			
Page 2:		1	1	1	1			
Page 3:			2	2	2*			
Hit/Fault	F	F	F	H	F			

LRU Example

Needed Page:	0	1	2	1	4	2	3	7
Page 1:	0*	0*	0*	0*	4	4		
Page 2:		1	1	1	1	1*		
Page 3:			2	2	2*	2		
Hit/Fault	F	F	F	H	F	H		

Needed Page:	0	1	2	1	4	2	3	7
Page 1:	0*	0*	0*	0*	4	4	4*	
Page 2:		1	1	1	1	1*	3	
Page 3:			2	2	2*	2	2	
Hit/Fault	F	F	F	H	F	H	F	

LRU Example

Needed Page:	0	1	2	1	4	2	3	7
Page 1:	0*	0*	0*	0*	4	4	4*	7
Page 2:		1	1	1	1	1*	3	3
Page 3:			2	2	2*	2	2	2*
Hit/Fault	F	F	F	H	F	H	F	F

$$p = \frac{6}{8} = .75$$

Note that LRU frequently (but not always) performs better than FIFO. There can and will be examples where FIFO will outperform LRU (Free Lunch Theorem)

- Replace the page in memory that will not be used for the longest amount of time
- Requires the entire stream of needed pages
 - Unless you have an amazing branch predictor, this will never work all of the time

Needed Page:	0	1	2	1	4	2	3	7
Page 1:								
Page 2:								
Page 3:								
Hit/Fault								

Optimal Example

Needed Page:	0	1	2	1	4	2	3	7
Page 1:	0(∞)							
Page 2:								
Page 3:								
Hit/Fault	F							

Needed Page:	0	1	2	1	4	2	3	7
Page 1:	0	0(∞) 1(3)						
Page 2:								
Page 3:								
Hit/Fault	F	F						

Optimal Example

Needed Page:	0	1	2	1	4	2	3	7
Page 1:	0	0	0(∞)					
Page 2:		1	1(3)					
Page 3:			2(5)					
Hit/Fault	F	F	F					

Needed Page:	0	1	2	1	4	2	3	7
Page 1:	0	0	0	0(∞)				
Page 2:		1	1	1(∞)				
Page 3:			2	2(5)				
Hit/Fault	F	F	F	H				

Optimal Example

Needed Page:	0	1	2	1	4	2	3	7
Page 1:	0	0	0	0	4(∞)			
Page 2:		1	1	1	1(∞)			
Page 3:			2	2	2(5)			
Hit/Fault	F	F	F	H	F			

Needed Page:	0	1	2	1	4	2	3	7
Page 1:	0	0	0	0	4	4(∞)		
Page 2:		1	1	1	1	1(∞)		
Page 3:			2	2	2	2(∞)		
Hit/Fault	F	F	F	H	F	H		

Optimal Example

Needed Page:	0	1	2	1	4	2	3	7
Page 1:	0	0	0	0	4	4	3(∞)	
Page 2:		1	1	1	1	1	1(∞)	
Page 3:			2	2	2	2	2(∞)	
Hit/Fault	F	F	F	H	F	H	F	

Needed Page:	0	1	2	1	4	2	3	7
Page 1:	0	0	0	0	4	4	3	3(∞)
Page 2:		1	1	1	1	1	1	7(∞)
Page 3:			2	2	2	2	2	2(∞)
Hit/Fault	F	F	F	H	F	H	F	F

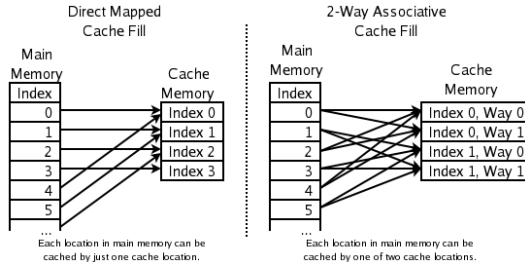
$$p = \frac{6}{8} = .75$$

- Cache is on chip memory used by the CPU to decrease the average time to access memory
 - Super fast super expensive memory
- Modern CPUs have 3 kinds of cache:
 - Data cache
 - Instruction Cache
 - Translation Look aside Buffer (TLB): Speeds up virtual to physical memory access for both instructions and data
- Cache hit: processor immediately reads data from or writes to the cache
- Cache miss: cache allocates a new entry reading in data from memory
 - Think page faults, just a lot less painful
- Cache lines: representation of a block of memory in cache

- Has 3 parts:
 - Tag: identifier for the cache line and part of the memory address
 - Data block: data for the cache line
 - flag bits: valid and dirty bits
 - valid: indicates whether the block has been loaded with valid data
 - dirty: indicates whether the block has been changed since it was read from memory.
Dirty means we need to save it to main memory.
- In general: If you hit the cache, the data lookup is free!

Associative Cache

- Decides where in cache a memory entry will go
- Fully associative: replacement policy may choose any entry
- Set association: memory entries are mapped to specific locations in cache
- Important in programming for High Performance systems: interleaving data in loops may allow for fewer cache misses



"Cache,associative-fill-both". Licensed under CC BY-SA 3.0 via Commons - <https://commons.wikimedia.org/wiki/File:Cache,associative-fill-both.png#/media/File:Cache,associative-fill-both.png>

Translation lookaside Buffer (TLB)

- Memory management hardware for cache
- Improves speed to translate virtual addresses
- Includes basic memory information such as references to pages, frames, and variables.

- Uniform Memory Access
 - Memory location accessed is independent of the processor
- Non-uniform memory access (NUMA)
 - Systems where memory access times vary wildly (based on memory location relative to the processor)
 - Memory location accessed depends on the processor
 - Every processor has it's own memory
 - Much like an extremely tightly coupled cluster on one machine
 - x86 Processors have supported NUMA since 2007

Cache Coherency

- Consistency of data on the cache shared by local CPUs.
- Three levels:
 - Write operations appear to occur instantly
 - Processor see the same changes of values for each operand
 - Different processor might see an op and assume different values (non-coherence)
- Three implementations
 - Snooping: individual caches monitor other caches
 - Snarfing: cache controller watches addresses and memory; updates its data copies once other data is updated in memory
 - Concurrency protocol: how do we implement coherence?
 - Hint: It's really hard
- ccnuma
 - allows for consistent memory image across sockets
 - IPC used to keep a consistent memory image when more than one cache stores the same memory location

- Virtual memory is constantly switching data between cache and memory, and between memory and disk
- Constant paging (memory to disk) can be resolved by:
 - MORE RAM!
 - Rewriting programs to be smarter
 - Reducing memory usage
 - Increasing spatial locality
- Spatial locality is part of the programmer's responsibilities

Example

- Download `examples/cache/cache_thrasher.c` from the course website
- Inspect the file and note the difference between the two functions being timed.
- Modify, compile, and run (cygwin should work for this) the program for the following values of `ARRAY_DIM`:
 - 10
 - 100
 - 1000
 - 5000
 - 10000
- Record the times and speedups that are reported in a table and answer the following questions:
 - What do you notice as the size of the 2D array grows?
 - Based on what you know about pointers and memory allocation and caching, what do you think is happening? Why is one faster than the other?